

# Hiring Assignment - Senior AI/ML Research Engineer

## Assessment: Research Innovation Lab

We have designed this assessment to evaluate your ability to approach complex problems in artificial intelligence research and implementation. The assessment consists of **two distinct challenges**, each targeting a specific area of expertise:

**Agentic Audio Intelligence for Cover Song Similarity** and **Title Normalization and Agentic Enrichment for Music Asset Titles**.

The candidate will perform a **full implementation of task A**, while **B** will be a **detailed description of the approach** that could be followed to address the problem.

## A. Implementation: Agentic Audio Intelligence for Cover Song Similarity

### Objective

This task evaluates your ability to design and implement an **AI agent-based system** that **quantifies and explains the similarity** between original songs and their cover versions using **multimodal data**. You are expected to leverage the **Da-TACOS dataset** and implement a **multi-tool agent** capable of reasoning over audio features, metadata, and (optionally) lyrics to produce a **similarity score** with **interpretable outputs**.

The goal is to demonstrate how **agentic systems** can enhance music similarity understanding in a modular and explainable way—critical for use cases in **digital rights management**, **copyright verification**, and **music recommendation**.

---

### Scenario

Imagine an AI Agent tasked with validating whether two musical recordings represent original/cover relationships. To perform this, it can activate a set of specialized tools that operate over different modalities—**audio**, **metadata**, and (optionally) **lyrics**—and then combine the outputs into a **final similarity score** with justification.

---

### Dataset

Use the [Da-TACOS Dataset](#), which includes:

- Cover cliques with known relationships
- Audio-derived features (Chroma, HPCP)
- Metadata (e.g., title, artist, release year)

You may optionally supplement with:

- Lyrics (via APIs or scraped metadata)
  - Pretrained music/audio embeddings (e.g., CLAP, MERT, MuLan)
- 

### Agent Architecture Requirements

You are required to build an **AI Agent** with at least the following **three tools**:

**Tool 1: Audio Embedding Similarity Tool**

- Uses pre-extracted audio features (e.g., chroma/HPCP from Da-TACOS) or pretrained audio embedding models (e.g., MERT, CLAP).
- Computes similarity between two tracks based on tonal, harmonic, or timbral alignment.
- Outputs: `audio_similarity_score`  $\in [0,1]$

**Tool 2: Metadata Comparison Tool**

- Parses metadata fields (e.g., title, artist name, year) and computes overlap/similarity.
- Can apply fuzzy string matching or embedding-based approaches (e.g., sentence transformers).
- Outputs: `metadata_similarity_score`  $\in [0,1]$  and justification (e.g., “same label”, “name overlap”)

**Tool 3: Multimodal Aggregator + Explanation Tool (Optional)**

- Combines scores from audio and metadata tools using a weighted scheme or a learned model (e.g., XGBoost, shallow MLP).
- Produces final similarity score along with a **natural language explanation** (e.g., “High tonal overlap and same release label detected”).
- Outputs:
  - `final_similarity_score`  $\in [0,1]$
  - `explanation: str`

You may optionally implement:

- A **Lyrics Embedding Tool** that extracts and compares semantic similarity using models like BERT or MPNet.
- A **Visual Explanation Tool** that uses SHAP/LIME to highlight input contributions.

---

**What to Build or Propose**

- A functioning AI Agent (script or framework) that takes two tracks (or IDs from the Da-TACOS dataset) and:
  - Invokes the tools
  - Aggregates results
  - Outputs a similarity score and human-readable explanation

---

**Evaluation Criteria**

Aspect	Criteria
Multimodal Reasoning	Effective use of audio features + metadata (optional lyrics)
Tool Design	Modularity and clarity in how each tool works
Explainability	Quality and interpretability of the generated justifications

<b>Scalability</b>	Ability to extend to new tools or larger datasets
<b>Code Quality</b>	Maintainability, modularity, documentation
<b>Research Thinking</b>	Creativity, understanding of limitations, alignment with current trends

---

### Deliverables

- A Python-based prototype implementing the agent and tools (script or notebook)
- CLI or function to run: `compare_tracks(song_id_1, song_id_2)`
- Sample outputs and evaluations (e.g., comparisons between known original/cover and unrelated songs)
- A short report (~2–4 pages) describing:
  - Agent architecture
  - Tools and models used
  - Design decisions and trade-offs
  - Limitations and ideas for future improvements

---

### Submission

- Upload your implementation to a public or private GitHub repository
- Include:
  - `README.md` with setup and usage instructions
  - Annotated code for tools and agent logic
  - PDF report or `REPORT.md`

---

### Encouraged Libraries & Tools

- **Audio:** Librosa, Essentia, TorchAudio, CLAP, MERT
- **Embeddings:** Sentence Transformers, BERT, MPNet
- **Similarity:** Scikit-learn, FAISS, cosine similarity
- **Explainability:** SHAP, LIME, attention visualization
- **Agents:** LangChain, OpenAgents, CrewAI (or simple Python logic mimicking agent behavior)

## B. Mini Survey: Title Normalization and Agentic Enrichment for Music Asset Titles

### Objective

This task evaluates your ability to design a **machine learning or LLM-based system** to **clean and normalize music asset titles**, while integrating **agentic reasoning**, **external knowledge enrichment**, and **explanation of the decision process**.

Assuming you have a dataset of annotated music asset titles, where each token has been labeled with categories such as clean title, artist name, version info, album title, or miscellaneous. Your job is to **design a comprehensive technical report** that showcases your ability to reason over noisy metadata and propose generalizable solutions with explainability.

---

### Dataset

**Asset Title Annotations** — A CSV dataset with annotated asset titles.

Each row contains:

- Full `raw_title` string
- Annotations indicating:
  - `title` (main title of the asset)
  - `artist` (e.g., performer or featured names)
  - `version` (e.g., Acoustic, Radio Edit)
  - `album_title` (if mentioned)
  - `misc` (e.g., country names, references, dates)

The dataset exhibits:

- Redundancies and false matches in version/artist segments
  - Sparse frequency for some tags like `album_title`, `misc`
- 

### Goal

Create or propose a system that:

1. **Cleans and normalizes the title** by isolating the true `title` field.
  2. Enriches ambiguous/missing parts using **external knowledge** (e.g., from Wikipedia, Spotify, or MusicBrainz).
  3. Produces a **transparent explanation** (via rules, XAI, or LLM reasoning) for how each segment of the input was classified or discarded.
- 

### Mini Research Survey and Technical Proposal

Deliver a detailed technical report that:

- Explores and compares possible modeling options
  - Proposes a modular agentic architecture for this task
  - Compares tools for enrichment (Spotify API, Wikipedia scraping, etc.)
  - Discusses evaluation metrics and error analysis strategies
  - Includes examples and architecture diagrams
- 

### Evaluation Criteria

Category	Description
----------	-------------

<b>Title Normalization</b>	Ability to identify and clean titles accurately using NLP/ML tools
<b>External Knowledge Use</b>	Ability to propose or use enrichment for disambiguation and validation
<b>Explanation &amp; Reasoning</b>	Use of interpretable outputs or explainable models
<b>Modularity</b>	Clear separation of components (classifier, enrichment, resolver)
<b>Research Skills</b>	Quality of the mini-survey or report (clarity, citations, comparison)
<b>Technical Depth</b>	Quality and creativity of chosen models, frameworks, or tools

---

## Deliverables

### Mini Survey:

- PDF or Markdown technical report
- Architecture diagram (if proposing Agentic system)
- Real examples with theoretical enrichment/explanation
- Comparative analysis of existing methods (optional: XAI, LLMs)

---

## Submission

- PDF report
- Bonus: mention limitations and potential for automation